

Qorvo

Semiconductor Manufacturer

During the manufacturing of semiconductor wafers Qorvo tests many individual circuits dozens of times to ensure quality. This creates a lot of data that is hard to analyze, they need a machine learning pipeline to identify patterns and anomalies in this dataset.

Solution

Pre-processing

In the pipeline there are preprocessing options allowing the users to define whether they want to impute missing values, correlated tests, or even aggregate test fails into a single column

Model

After preprocessing the data is given to a VAE. The VAE, a variational autoencoder, begins by training on the data in “reverse” taking images, compressing them down, reconstructing the original image from the compressed data and comparing that to the original image, creating a loss value. This is repeated until the loss is at an acceptable level. The model is then run normally on the data, creating a similarity matrix by comparing the compressed data in the latent space.

Clustering

This similarity matrix is then translated to a distance equivalent:

$$Dm = \sqrt{2 * (1 - Sm)}$$

The distance matrix is then passed to either a hierarchical or k-means clustering algorithm depending on whether the user wanted a specific or dynamic number of groups. This creates groups of wafers based on the similarity created by the VAE.

Who are We?

We are a group of students at Oregon State University - Cascades in our senior capstone program. We were provided a number of projects proposed by local companies to work on as pseudo interns to get work experience. We ended up on this project for Qorvo. We found the data tough to start working with having vague labels, dozens of columns, and tens of thousands of rows. Once we understood the data we were able to start looking at models to use, it took a couple months, but eventually we found what we were looking for. With this we split into two groups, with one person working on fine tuning the model and understanding it, and the other two people creating pre and post processing steps.

- The technical details of your implementation. What is the tech stack? How does it all work? (Be brief.)

Environment

Databricks

The entire project is contained within databricks in a series of jupyter notebooks

Notebooks

We have two primary folders of jupyter notebooks in our directory, one where we did all of our development, and one where we run the pipeline on multiple datasets using notebooks for each with documentation for every step. Within this folder there is another that holds all the notebooks for the pipelines source code, containing files for every step with pydocs for every function.

- The current status of the codebase. Where did you stop? What do you recommend someone does next?

Current Status

Next Steps

Continuous classification

The model only handles the whole dataset provided to it right now, however it would be useful to introduce new data to an already trained model, returning a new matrix with this new data.

- Notes, highlights and concerns. What are some vague areas of the system? What needs attention? What should a new developer watch out for?
- How to get up and running in development, and how to build/deploy the system to the production environment. (This can be brief, if Part 2 below is sound.)

Deployment

Notebooks

A user is able to make a copy of one of the existing pipeline notebooks, change the database to the data they want to run the pipeline on, adjust settings on the steps used, and adjust the input size for the model. The user could also create their own from scratch and use the functions provided whoever they wish, skipping preprocessing steps as they wish.